

mHot

**Mobile Here or There - A framework for runtime
offloading of computation from mobile devices to the
cloud.**

by

William Joseph Gaudet

B.ScEng, The University of New Brunswick, 2008

A THESIS SUBMITTED IN PARTIAL FULFILLMENT OF
THE REQUIREMENTS FOR THE DEGREE OF

MASTERS OF APPLIED SCIENCE

in

The Faculty of Electrical and Computer Engineering

(Engineering)

THE UNIVERSITY OF BRITISH COLUMBIA

(Vancouver)

July 2011

© William Joseph Gaudet 2011

Abstract

Today the average north american carries a computer several orders of magnitude more powerful than the computer used to guide the Apollo space craft to the moon and safely back. However, with this dramatic increase in both ubiquity and power of personal mobile computing, has come a host of applications that require significantly more computational power than is available in today's smart phones. Additionally, battery life of mobile devices is still a significantly limiting factor when doing large quantities of computation on mobile platforms.

With these factors in mind this research investigates the gains in performance, quality of computation, and battery life which can be made possible by real time offloading of computation from a mobile device to a server platform.

Furthermore, it details the design and implementation of a framework for offloading computation from mobile devices which has at its core the goal of limiting the cognitive overhead for the developer to conduct such an offload.

Table of Contents

Abstract	ii
Table of Contents	iii
List of Tables	v
List of Figures	vi
List of Programs	vii
Acknowledgements	viii
Dedication	ix
1 Introduction	1
1.0.1 Decision Goals	2
2 Background and Previous Work	4
3 Design and Modelling	5
3.1 Overview	5
3.2 A Framework for Data Migration	5
3.2.1 Dynamic Proxies	5
3.2.2 Remote Compute Service	10
3.3 Programmer Responsibilities	11
3.4 The Model	11
3.5 Exercising the Framework	11
3.5.1 Basic Experiment	11
4 Experimental Results	12

Table of Contents

5 Conclusion	13
5.1 Summary of Contribution	13
5.2 Future Work	13

Appendices

First Appendix	14
Second Appendix	16

List of Tables

1 Table of field sizes 15

List of Figures

3.1	An Activity Diagram showing the RMI invocation process . .	7
3.2	The general form of serialized RemoteObjects	10

List of Programs

Acknowledgements

Dedication

Chapter 1

Introduction

Motivations

Desktop and laptop computers have been the mainstay of personal computing for more than 30 years, providing untold utility to millions of users worldwide. However, due to recent improvements in battery technology paired with reduction in size and cost of the components necessary to build a computer system, industry leaders are declaring the "Death of the PC" and subsequent rise of mobile platforms, referring to smartphones and tablet computers rather than laptops. At the All Things Digital Conference 2010 Steve Jobs, founder and CEO of Apple, echoed this sentiment, stating that Personal Computers (PCs) would become like trucks: they will still exist but the bulk of people won't need them.

Additionally, this new era of pervasive mobile computing has given rise to a suite of new mobile applications, with the average smartphone boasting the equivalent computational capability of a PC from 2001, the potential for innovative new applications seems limitless—this is aided by the many environmental sensors present on the typical smart phone such as those which measure position, orientation, and acceleration.

There is however a second class of application that requires significantly greater capabilities than are offered by the current generation of mobile platforms. An example of this class of application include music recognition programs such as Shazam [CITATION] which requires a large corpus of music against which to compare an audio sample, additionally the digital signal processing requirements are likely greater than the capabilities of the mobile device (or they process would be prohibitively non performant on the mobile device). In spite of these limitations, these applications exist and provide a great deal of utility to the end user. The shortcomings of the mobile devices are addressed through the offloading of the more challenging computational work to a computer with greater resources and capabilities at its disposal—This process will henceforth be referred to computational offloading. The benefits of computational offloading are seemingly obvious; providing the user the facade of having a more performant computer than they already

have. In fact computation offload is not a new idea, search systems such as Google can be viewed as computational offloading systems. Whereby the user provides a search string into the google web portal, which has significantly more capabilities and data at its disposal, and google returns a listing of search results.

Currently if a developer wants to add some computational offloading to their mobile application, they must explicitly construct their own method for offloading their data structures, additionally they must engineer a remote runtime to accept the offloaded data, execute the remote operations, and return the results to the mobile device. This gives rise to a large number of potential problems for a developer. There are many parameters which must be considered when offloading computation, not the least of which is the bandwidth of the remote connection. In a mobile scenario there is a delicate balance between the bit rate of the connection and the power consumption costs (watts/byte) which must be considered when determining the value of a computation offload. Often times these complexities give rise to incomplete or inefficient solutions, as the developer may lack the experience and skill to create a reasonable offloading system.

The goal of this work is two fold, to identify the critical parameters which must be taken into account in order to develop a decision model for offloading computation, and to describe and implement a framework for the android mobile platform which limits the cognitive overhead for the developer when creating mobile applications which require computation offloading.

1.0.1 Decision Goals

The decision model described in this work will focus around three goals, minimizing runtime, maximizing computation 'quality', and maximizing battery life. Minimizing runtime is a fairly straightforward goal, the decision to offload computation can be made trivially by predicting the which platform will have the lowest runtime (the following equation is used).

$$\min(T_{compRemote} + T_{transmission}, T_{compLocal}) \quad (1.1)$$

Computation quality is considered in a special case of problem, where an increased number of iterations of a particular problem would increase the correctness of a solution. An example of this would be an iterative root finding method, where more iterations would lead to a more correct result (a value closer to the actual root). In the case of a mobile computation offloading situation, one could reason that if a problem were executed in the same approximate running time but on server with increased capabilities the

over all quality of the computation could be increased. Or put a different way, given a boundary of some level of quality of result, the running time could be minimized.

$$\min(T_{compRemoteToQuality(q)} + T_{transmission}, T_{compLocalToQuality(q)}) \quad (1.2)$$

Finally in general the goal of reducing battery life can be expressed simply as:

$$\min(P_{compLocal}, P_{transmission}) \quad (1.3)$$

where P is the cost in watts of a particular operation.

Development Overhead

Minimizing the cognitive overhead of developers is the goal of most development frameworks, reducing the amount of work a developer has to do both in lines of code and in adjusting their mental picture of a problem in order to solve it is in general viewed as desirable. The design of this framework has limited the amount of input from the developer to a sole annotation on a method, the use of some custom data structures when building their domain objects, and the following of some development best practices. These will all be described at length in the design and implementation section of this document.

Chapter 2

Background and Previous Work

As was stated in the introduction, computational offload is not a novel concept per say. There are several applications available to smart phone users today that make use of computational offloading to provide a more rich experience to the end user. Shazam was one such application, where a user records a section of music and the app sends the recording to a server for identification. [ADD SOME MORE EXAMPLES]

Cuervo et al [?], presented their framework MAUI, which allowed for runtime offloading of mobile computation from the windows platform. The MAUI framework sought to maximize battery life by determining at runtime which problems could be executed remotely on a parallel runtime. However their approach had some flaws, their decision to allow the developer to use regular primitive data types and rely on reflection and serialization to ship their objects to the remote server while desirable for the developer in terms of ease of use, the use of reflection and serialization incur processor overheads which contribute to the power cost of the object transmission.

Chapter 3

Design and Modelling

3.1 Overview

3.2 A Framework for Data Migration

3.2.1 Dynamic Proxies

The primary mechanism for abstracting the decision of where to execute a method (in the cloud or on the mobile device) is the Java Dynamic Proxy class [citation]. The dynamic proxy class enables runtime proxying of any object provided it implements an interface. A dynamic proxy object then implements the interface of the object to be proxied at runtime, and then routes all method calls through its `invoke` method. It is in this `invoke` method that the decision to either package the object graph and ship it to the server for execution or execute locally is made.

The dynamic proxy class does have some shortcomings, it requires developers to create interfaces for the methods they want to be remotable, additionally the developer must explicitly create the proxy objects for use in their code, and finally dynamic proxies do incur a performance overhead. Simple running time experiments show that the average time cost of a regular Java method invocation is approximately 60ns while the average time cost of the same method invocation done through a dynamic proxy object is around 5ms. However with these costs in mind a developer can make architectural decisions to help mitigate these performance overheads. The shortcomings with respect to the use of interfaces is actually viewed as a positive as it forces the developer towards best practices of programming to interfaces and not implementation. Finally steps have been taken to mitigate developer pain with regards to creation and use of dynamic proxies, by way of the Java Generics system every object which descends from the `AbstractRemoteObject` has a factory method which will wrap the object in a proxy.

Remote Method Invocation

When mHot determines that a method should be executed in the cloud rather than locally on the mobile device, it must conduct a remote method invocation. While it takes its name from the Java RMI system, the remote method invocation object is only related to RMI in the fact that it allows for a developer to remotely invoke methods on an object. The Java RMI system is a fully feature interprocess communication system and as such when mHot was first designed RMI was investigated as a potential solution to the problem of object transmission and method invocation. However, the lack of support for native RMI in the Android platform was seen as an insurmountable barrier. As such, a custom and lightweight remote invocation methodology was devised with a goal in mind to minimize both the amount of computation required to marshal the objects between the compute platforms and limit the amount of data that would be transmitted over the wire—this is particularly important due to battery life constraints and the substantial energy costs of the high bandwidth radios (3g / WiFi).

A single remote invocation is a collaboration between three objects, the target remotable object, the argument array, and the result. Each object plays a particular role in the method invocation process, and together make up the composition of a `RemoteInvocationMethod` object. The target remotable object is the object against which the method will be invoked, the argument array contains facilities for serializing and deserializing the arguments the method takes—these objects must be either primitive data types, or implement `RemoteableObject` otherwise an exception will be thrown—, finally the result is container for the returned result of the method. The remote method invocation process is performed by the Remote Dynamic Proxy Object, this process is outlined in the activity diagram shown in Figure 3.1. Presently the remotated method should not mutate the target object’s state as mHot has no facilities for returning the object once it’s been shipped to the remote server, thus the developer must currently take that into account when designing their APIs. However this could reasonably be addressed via some post execution synchronization mechanism, or by returning the mutated object as the result—though this would be a poor stop gap solution as it would increase the amount of traffic on the line.

Data Serialization

When a dynamic proxy object receives a method invocation that has been tagged as remotable, if the framework has determined that the request

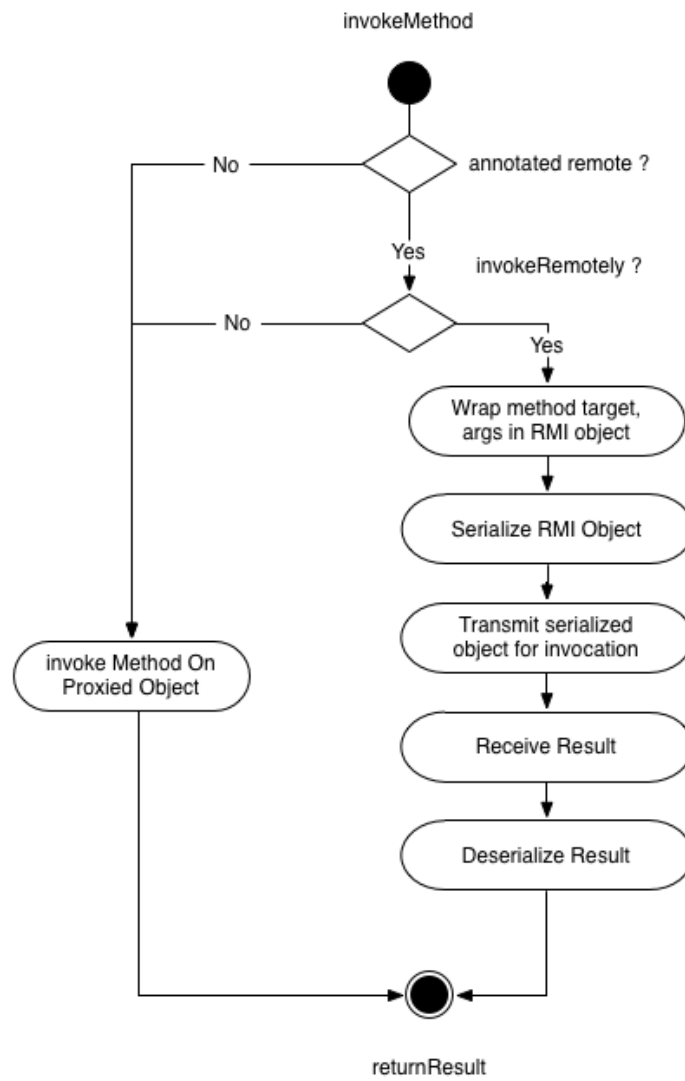


Figure 3.1: An Activity Diagram showing the RMI invocation process

should be executed remotely, the proxied object must first be serialized before it can be shipped to the server. The proxied object can take on one of two forms: a simple object composed only of primitive data types (integers, floating point numbers, booleans etc) or a complex object graph composed of zero or more primitive types and one or more references to other objects. These other objects may in turn also contain references to additional objects both complex and simple as well as cyclical references between objects.

Given the cost associated with object serialization both in terms of computation (due to memory allocation and object graph traversal) and in terms of communications bandwidth the task of the object serializer becomes two fold. Most importantly it must correctly serialize the object graph preserving any and all object references (including circular ones); additionally the serializer must attempt to minimize the size of the transmitted object graph.

This is primarily done by the encapsulation of an object's fields in container object which implements the `RemotableField<T>` interface. [CODE LISTING] The interface exposes several values about the encapsulated object all of which are required to correctly and efficiently (in both space and time) serialize the object graph, the more relevant methods to the serialization process are the following:

- `int size()` - Returns the size in bytes required to serialize this field (in the implemented case of a `RemoteFieldObject` this is 4 bytes for the reference to the object)
- `<T> get()` - Returns the object which is encapsulated in this field, in the case of a primitive value it is wrapped in it's Java object container type (`Integer` for `int`) etc, due to performance issue with the autoboxing system [CITATION NEEDED] where `T` is either a primitive data type or an implementation of `RemoteObject`
- `boolean isDirty()` - Returns true if this object has had it's value set since initialization or last serialization - this is used to minimize the amount of data that must be transmitted to the remote server.
- `void Serialize(ByteBuffer)` - Serializes the field to the provided `ByteBuffer`
- `void Deserialize(ByteBuffer)` - Deserializes the field from the provided `ByteBuffer`
- `FieldVisitor acceptVisitor(FieldVisitor)` - Accepts the field visitors which are used to conduct a variety of tasks in the graph serialization process.

The Serialization Process

The object serialization process is a modified depth first search. In effect the process visits every node in the graph twice, once to compute the overall size required to encode the graph as a binary representation, and a second time to perform the serialization it self. [This is a potential future work situation, some of this could probably be improved].

When a object in the graph is visited by the size computation pass, each of the object's remotable fields are interrogated for their serialization size—It should be noted that the fields for each object class are stored in a memoization cache to prevent recalculation of the classes' remotable fields through the use of the Java reflection APIs which are not very performant [Maybe this should be a footnote probably needs a citation]. If the field encapsulates a primitive field it simply looks up the serialization size cost for that field which is its data size in addition to the field reference, however, if the field encapsulates a remotable object it either adds 8 bytes if it has been visited once before—the byte cost of referencing an object which has already been visited—or it adds 8 bytes plus the serialization cost of that object, which is computed immediately—A complete list of default data type serialization costs is shown in 1. The size cost of an object is computed by adding the length of the fully qualified class name (eq: com.someCompany.somePackage.someClass) to the size of the serialized data including referenced object sizes to 8 bytes—4 for the entire schema size, 4 for the first object size.

On the second traversal of the object graph each object is serialized by calling each of its fields serialize method. This method creates a binary representation of the data which can later be decoded to reproduce the same values—The put and get methods on the Java ByteBuffer class handle the serialization of primitive data types into their byte representations. A primitive field is serialized as a tuple of the hash value of the string representation of it's name and it's value. The hash value is used to reduce the size and complexity of the encoding. Given that a unique string of arbitrary length will have a unique hash code, this allows each field to be key value encoded with a 4 byte integer for the code. As such a single field is be represented by a 3 tuple of the length of the field representation (including the length of the length of the length which is a 4 byte integer), the hash code of the field name, and the value of the field. ObjectFields—that is fields which reference remote objects—are represented as a length, field name hash code, and the hash code of the object they reference, the object it self is added to a queue of objects to be serialized. The general form of a serialized remote object

3.2. A Framework for Data Migration




Figure 3.2: The general form of serialized RemoteObjects

post serialization has completed is shown in Figure 3.2

3.2.2 Remote Compute Service

The remote compute service (RCS) is a simple socket server and key value store, that runs on the cloud compute platform. It's primary task is to listen for incoming invocation requests, upon receipt of a request it deserializes the target object and the argument array. These deserialized objects are synchronized with the local cache of remoted objects—Objects are sent as deltas only to prevent retransmission of information. Once the cache is synchronized the method is invoked against the target object with the argument array, and the result is in turn serialized and transmitted to the mobile device.

Deserialization is relatively straightforward, however it requires the RCS to have access to all of the class files that represent the objects that have been transmitted. This is done by simply jarring all pertinent files and having them present in the /jars folder which is located at the root of the classpath of the RCS. A single object is deserialized by loading it's class by name, and reflectively instantiating it. This requires that all RemoteObject have a null constructor (an constructor with no arguments). After which fields are looked up by the hash value of their field name and set to the serialized value.

3.3 Programmer Responsibilities

3.4 The Model

3.5 Exercising the Framework

3.5.1 Basic Experiment

The first and most simple experiment will be used to determine if it is possible to obtain more than marginal gains in terms of battery life and computation time (over all number of computations executed in the given battery life) using a non heuristically or stochastically based model, but rather by simply shipping the computation off to a server with greater capabilities.

Chapter 4

Experimental Results

Chapter 5

Conclusion

5.1 Summary of Contribution

5.2 Future Work

First Appendix

Type	Data Size	Effect On Schema
byte or Byte	1	5
char or Character	2	6
short or Short	2	6
float, Float, int, or Integer	4	8
double, Double, long, or Long	8	12
byte[length]	1 * length	1 * length + 4
char[length]	2 * length	2 * length + 4
float[length] or int[length]	4 * length	4 * length + 4
double[length] or long[length]	8 * length	8 * length + 4
object (not visited)	object.size()	object.size() + 8
object (visited)	object.size()	8

Table 1: Table of field sizes

Second Appendix

Here is the second appendix.